BLADE-NCO Theory Research
Stages 1–3B: Reachability, Hardness Tracking, and Identifiability

Research note

October 22, 2026

Abstract

This note formalizes structural reachability, inherited population memory, moving hard-instance regions, and first discovery measured in computational cost. It records three candidate claims, attempts a proof of a finite-window memory comparison, and tests broad memory/diversity claims against exact finite-state examples. The strongest surviving result is conditional: retained structural memory helps only when it supplies a verified per-cost hit-hazard advantage over the comparator. No unconditional BLADE-NCO advantage is established. This document is a cumulative Stages 1–3B mathematical research note, not a completed theory paper.

Stage 1.5 corrects the state space, time indices, probability space, aggregation language, and implementation mapping without adding a theorem or changing that conclusion. The finite examples remain counterexamples to universal claims; they are not assumptions that the TSP instance space is finite.

Stage 2 proves a deterministic occupied-niche persistence invariant for the verified parent–child union and Elite-QD selector, while showing by exact counterexample that niche persistence need not preserve useful structural representatives. It also gives a finite-budget multistep bridge-chain bound whose factors are operator-specific accepted-transition and selection-retention certificates. The theorem is mathematically valid under its stated premises; the decisive BLADE-specific kernel and bridge-retention premises remain open.

Stage 2.1 independently corrects the information set for deterministic versus probabilistic bridge retention, imposes a non-anticipating choice of cycle-indexed bridge targets, makes hash-deduplication compatibility explicit, and separates hard-instance presence from new-instance discovery. These are corrections to the conditional framework, not additional superiority theorems.

Stage 3A proves a standard per-niche approximate-argmax guarantee and its one-step solver-drift extension. It also derives a finite-horizon historical tracking-regret recurrence under a substantially stronger drift bound over the complete latent candidate history. The recurrence can accumulate linearly, and its complete-pool HGB-error and historical-drift premises are not verified for S6-AE. Exact counterexamples continue to separate current-hardness tracking from future reachability and downstream BQ training value.

Stage 3B removes those unsupported premises and audits what is observable. It proves an exact exchange identity between Elite-QD acquisition and pure global HGB Top10k, making the surrogate-score opportunity cost computable. An exact finite-pool indistinguishability construction shows that selectively verified labels do not identify full-pool niche maxima, elite true-hardness regret, hard-tail recall, or the true Elite-QD-versus-global comparison. The missing paired comparison can be completed with at most 1,000 same-checkpoint labels, but those labels are not present in the frozen artifacts.

1 Scope and status

The scope is restricted to structural reachability, evolutionary memory, and expected first-hitting time under a non-stationary hardness landscape. The supplied BLADE-NCO context is FOGA-11 structural mutation, inherited parent populations, HGB40 hardness estimation, Elite-QD1000 (P=T), and BQ solver co-adaptation. Those names are treated as labels, not as complete algorithm specifications.

Statuses are reported on three separate axes. [**Mathematical status: PROVED**], for example, concerns the validity of a mathematical statement under its premises; [**Implementation applicability: VERIFIED, PARTIAL, OPEN, or NOT APPLICABLE**] concerns whether those premises and variables match the inspected S6-AE implementation; and [**Scientific novelty: OPEN or NOT CLAIMED**] concerns literature priority. A proof of a conditional probability bound does not establish implementation applicability or novelty. Compilation verifies only that the source typesets; it does not verify mathematical correctness.

The Stage 1.5–2 implementation boundary is the verified HGB40 + Elite-QD1000 ($P = T$) S6-AE path. Its frozen artifacts passed 40/40 strict checks. Historical GP4 output is not used as the definition of the current method. Details that could not be established from the inspected code and artifacts are marked OPEN rather than inferred.

2 Prior-theory boundary

Related theory includes rigorous results in which populations track moving optima and in which diversity mechanisms have different effects, but those works analyze explicit dynamic-function classes and specified algorithm dynamics [1, 2]. MAP-Elites and QD explicitly store high-quality representatives in a fixed feature partition [3, 4]; recent runtime analyses prove coverage and optimization results for particular pseudo-Boolean and combinatorial structures [5, 6]. These works make the elementary archive-occupancy invariant non-novel, and they show that stepping-stone benefits require an aligned problem/operator model. They do not establish compute-normalized discovery of solver-relative hard TSP instances under surrogate selection and solver co-adaptation. Whether a fully instantiated BLADE-specific bridge theorem would be novel remains [**Scientific novelty: OPEN**]; the present limited audit is not a comprehensive priority search.

3 Formal model and implementation boundary

3.1 Structural states and mutation

Let $\mathcal{Z}$ be the set of encoded candidate instances and let $\phi : \mathcal{Z} \to \mathcal{X}$ be a measurable map from encodings to the solver-visible TSP instance. For the inspected fixed-size experiments one may take $\mathcal{X}$ to be the valid canonical subset of $[0, 1]^{100\times 2}$, equipped with its Borel σ-algebra $\mathcal{B}_{\mathcal{X}}$; a variable-size model may instead use a measurable disjoint union over instance sizes. Encodings z, z' are equivalent when $\phi(z) = \phi(z')$. A structural state is a phenotype $x \in \mathcal{X}$, for example a canonical graph/instance together with whichever structural attributes determine solver input. The exact canonicalization and attributes are implementation questions. We assume ϕ and every event used

below are measurable.

The instance space is not made finite by a finite search budget. A realized finite-budget run contains finitely many attempted instances, but the support of a continuous-coordinate mutation kernel can remain uncountable. For each BQ cycle k and FOGA operator o, adjoin an isolated rejection symbol ∂ and write $\overline{\mathcal{X}} = \mathcal{X} \dot{\cup} \{\partial\}$ with the natural extended sigma-algebra. Let

$$\widetilde{K}_{k,o}(x, \cdot \mid s)$$

be a Markov kernel from a parent x and full pre-mutation state s into $\overline{\mathcal{X}}$. An outcome in $\mathcal{X}$ is the valid repaired phenotype of an accepted proposal; ∂ records a rejected or invalid proposal. The kernel includes stochastic operator parameters, repair, and canonicalization. Dependence on k or s permits solver-, archive-, or history-dependent behavior. The corresponding encoding process need not remain Markov after projection through ϕ.

Let $d : \mathcal{X} \to \mathbb{R}^4$ be the measurable QD descriptor map, including the frozen normalization, and let $g : \mathcal{X} \to \{1, \ldots, K\}$ be nearest-centroid assignment with a declared measurable tie rule. Then $\mathcal{N}_j = g^{-1}(\{j\})$ gives a finite measurable partition after discarding any empty cells. Finiteness belongs to the niche index set, not to $\mathcal{X}$.

3.2 Two time scales and the S6-AE batch

Use k only for the BQ adaptation cycle. At the start of cycle k, the BQ solver state q_k and the parent population $P_k = (X_{k,1}, \ldots, X_{k,m})$ are fixed, with $m = 10{,}000$ in S6-AE. Let P_k also denote its membership set when used inside $\subseteq$, $\cap$, or $\mathsf{S}(\cdot)$; the ordered tuple is used for scheduling. Let $\mathcal{O}$ be the ordered list of the 11 FOGA operators. The within-cycle mutation attempts are indexed by $(i, o) \in \{1, \ldots, m\} \times \mathcal{O}$:

$$Z_{k,i,o} \sim \widetilde{K}_{k,o}(X_{k,i}, \cdot \mid S_k).$$

This notation denotes a conditional law, not random parent or operator selection. The inspected implementation uses the deterministic Cartesian schedule "parents in stored order, then operators in registry order," so each parent–operator pair is attempted exactly once. A content-derived seed from $(k, \mathrm{hash}(X_{k,i}), o)$ drives the stochastic operator parameters and repair. If $Z_{k,i,o} = \partial$, that scheduled attempt produces no child and is not repeated. Conditioning the kernel on acceptance would instead describe an implicit-resampling algorithm and is not used here.

Let FU retain first occurrences in a finite ordered list by instance hash. Apply it to the subsequence of $(Z_{k,i,o})_{i,o}$ lying in $\mathcal{X}$ to obtain the accepted, child-deduplicated list O_k. An accepted duplicate is removed by this deterministic map; it is not redrawn. Let $C_k = \mathrm{unique}(P_k \cup O_k)$ be the candidate union, retaining the first occurrence under the declared order. In statements that treat hash-based deduplication as mathematical set union, assume that equal hashes identify the same solver-visible phenotype on the realized finite pool. For the weaker occupied-niche equality alone, it is sufficient that equal-hash candidates have the same niche label:

$$\mathrm{hash}(x) = \mathrm{hash}(y) \implies g(x) = g(y) \qquad (x, y \text{ in the pre-deduplication pool}).$$

Without this compatibility assumption, first-occurrence deduplication could erase a genuinely different offspring niche. The implementation checks content-derived hashes operationally; cryptographic collision resistance is not a literal mathematical injectivity theorem. At C00 S6-AE reuses the hash-locked corrected union; later cycles construct this union directly. The cycle then fits/scores

with HGB40, selects the acquisition/next-parent set, obtains exact labels for that selected set, and only then adapts the BQ solver from q_k to q_{k+1}. Thus within-cycle mutation attempts and solver cycles must not share one index.

3.3 Population memory and parent survival

At BQ cycle k, let P_k be the ordered parent population and A_k the branch-local labeled archive. Let S_k collect scores, lineage information, solver state, and other algorithm variables that affect the next transition. Thus $M_k = (P_k, A_k, S_k)$ is the operational memory state. Parent survival is a transition event: for parent i, its survival probability is

$$r_{k,i} = \mathbb{P}(X_{k,i} \in P_{k+1} \mid \mathcal{F}_k^{\mathrm{pre}}).$$

In a generic policy this may differ from a proposal-selection probability. For S6-AE, however, no such random parent-selection variable is present: every member of P_k receives all 11 scheduled mutations. Population update is a deterministic function of the realized candidate union, HGB40 scores, niche assignments, and hash tie-breaks. Randomness can still enter through the roots, mutation parameters, calibration sample, model fitting, training permutations, or nondeterministic numerical execution. The full process may be time-inhomogeneous Markovian in a sufficiently rich (M_k, θ_k, q_k); a niche label or count vector alone need not be.

For a finite candidate set A, define its occupied-niche set by

$$\mathsf{S}(A) = \{g(x) : x \in A\}.$$

For a finite unique candidate union C and score $\widehat{h}$, let $e_j(C)$ be the maximum-score member of $C \cap \mathcal{N}_j$, using ascending instance hash to break ties, and define the per-niche floor

$$Q(C) = \{e_j(C) : j \in \mathsf{S}(C)\}.$$

This set has $|\mathsf{S}(C)|$ members, not 10,000 members in general. If the population capacity is m, $|\mathsf{S}(C)| \le m \le |C|$, let $F_m(C)$ be the first $m - |Q(C)|$ candidates of $C \setminus Q(C)$ under global descending $\widehat{h}$ and ascending-hash tie order. The complete selected population is

$$T_m(C) = Q(C) \cup F_m(C).$$

The notation denotes the membership set; the implementation also records a deterministic order.

In verified S6-AE, $m = 10{,}000$, $K = 1000$, $Q_k = Q(C_k)$, and $T_k = T_m(C_k)$. The result has 10,000 unique members and is used both as the BQ acquisition set and next parent population: $P_{k+1} = T_k$. "Elite-QD1000" therefore names the occupied-niche floor Q_k, not the full T_k. The score can change at every cycle without changing the fixed partition g.

For any $x \in C$, the deterministic membership criterion is

$$x \in T_m(C) \quad \Longleftrightarrow \quad x = e_{g(x)}(C) \text{ or } \left[x \notin Q(C) \text{ and } \operatorname{rank}_{C \setminus Q(C)}(x) \le m - |Q(C)|\right].$$

Consequently, retaining a niche does not imply retaining any particular parent hash in that niche.

3.4 Causal information structure and time-indexed bridge sets

For each cycle t, let $\mathcal{F}_t^{\mathrm{pre}}$ represent all information available immediately before resolving the designated mutation batch: the entering solver state q_t, available parents and archive, and any previously realized random inputs and scores. It must exclude randomness of the batch still to be drawn under the probability ensemble. Let $\mathcal{F}_t^{\mathrm{batch}}$ add the realized accepted/rejected outputs and candidate union, and let $\mathcal{F}_t^{\mathrm{score}}$ additionally contain all scores, niche assignments, and tie-break data necessary for selection. With the verified deterministic Elite-QD rule, the event that a particular bridge is retained is measurable with respect to $\mathcal{F}_t^{\mathrm{score}}$, and its conditional probability given that complete information is its indicator (zero or one). The probabilistic retention premise below is therefore conditioned on the *pre-batch* information, not the fully resolved selection state.

For a prospective chain starting at cycle k, choose measurable bridge sets $B_0, \dots, B_L$ causally. The initial set B_0 is known at $\mathcal{F}_k^{\mathrm{pre}}$. For $\ell = 0, \dots, L-1$, choose the next set $B_{\ell+1}$ using only $\mathcal{F}_{k+\ell}^{\mathrm{pre}}$, *before* drawing that cycle's batch, and keep its identity fixed until the next retention step. A measurable random-set formulation is allowed if its membership event is jointly measurable. Thus B_ℓ is already fixed when it is used as a parent-availability certificate at cycle $k+\ell$. No bridge set may be chosen retroactively to fit a successful offspring.

Let

$$E_\ell = \{P_{k+\ell} \cap B_\ell \neq \varnothing\}, \qquad G_\ell = \{C_{k+\ell} \cap B_{\ell+1} \neq \varnothing\},$$

and write

$$R_\ell = \{T_{k+\ell} \cap B_{\ell+1} \neq \varnothing\}.$$

Since $P_{k+\ell+1} = T_{k+\ell}$, $R_\ell = E_{\ell+1}$. E_ℓ is measurable before cycle $k+\ell$'s generation, whereas G_ℓ and R_ℓ become measurable after generation and selection. The bridge-success prefix is $H_0 = E_0$ and $H_{\ell+1} = H_\ell \cap G_\ell \cap R_\ell$.

For terminal correctness, the final B_L must be selected at or before $\mathcal{F}_{k+L-1}^{\mathrm{pre}}$, when the entering solver q_{k+L-1} is known, and satisfy

$$B_L \subseteq \mathcal{H}_{k+L-1} \quad \text{almost surely on admissible pre-generation histories.}$$

This condition allows the final target to depend causally on the current solver, but forbids selection of B_L using an unseen future checkpoint or the realized final offspring. If all bridge sets are instead frozen at cycle k, terminal inclusion must hold across the subsequently possible solver trajectories; this is a stronger, robust-target assumption.

A prospective future-reachability value can be written

$$V_{k,L}(x) = \mathbb{P}(\text{a declared future true-hard set is hit within } L \text{ cycles} \mid x \text{ is retained at cycle } k,\ \mathcal{F}_k).$$

This value generally depends on other population members, future selection, and the adapted solver states. It need not be determined by current BQ hardness or current HGB prediction.

3.5 Probability space and frozen executions

Let $(\Omega, \mathcal{F}, \mathbb{P})$ carry exactly the random inputs declared for the analysis: possibly random initial roots, content-derived operator seeds and their stochastic parameters/repair, seeded archive calibration sampling, HGB estimator seeds, BQ training permutations and suffix sampling, and any

explicit selection randomness in a comparator. Conditioning on the frozen S6-AE roots, source artifacts, software/runtime, and all recorded seeds gives a single realized trajectory ω_0; many nominal "probabilities" then collapse to zero or one. A probabilistic reachability theorem must therefore state which inputs vary across Ω. Repeatedly serializing a deterministic batch does not create random parent/operator selection, and unspecified hardware nondeterminism is not silently treated as a scientific sampling law.

3.6 Hardness and moving target

Let θ_k describe the task/environment and q_k the state of the co-adapting BQ solver before cycle-k adaptation. Define true BQ hardness $h_k^{\mathrm{BQ}}(x) = h(x; \theta_k, q_k)$, where a larger value means harder under a prespecified solver-performance metric. For threshold κ, the true hard region at cycle k is

$$\mathcal{H}_k = \{x \in \mathcal{X} : h_k^{\mathrm{BQ}}(x) \geq \kappa\}.$$

This definition allows the target to move because either the task or solver state changes. During the S6-AE mutation/acquisition part of cycle k, q_k is fixed; adaptation produces q_{k+1} only after T_k is frozen and labeled. For the proofs, the target attached to a declared opportunity must be predictable from its pre-opportunity history.

HGB40 supplies a cycle-specific prediction $\widehat{h}_k(x)$ (in the implementation, predicted $\log(1 + \mathrm{LBQ})$), not the definition of h_k^{BQ} or $\mathcal{H}_k$. Define the latent generation-level event

$$D_k^{\mathrm{gen}} = \{\exists x \in C_k : h_k^{\mathrm{BQ}}(x) \geq \kappa\},$$

and the acquisition-level event

$$D_k^{\mathrm{acq}} = \{\exists x \in T_k : h_k^{\mathrm{BQ}}(x) \geq \kappa\}.$$

The former concerns what the generation-level candidate union contains, including retained parents and accepted children, even when exact hardness was not measured; the latter concerns what survived HGB/QD acquisition and can be verified. They are distinct from the event $\{\widehat{h}_k(x) \geq \widehat{\kappa}\}$ and from post-adaptation hardness $h_{k+1}^{\mathrm{BQ}}(x)$. Without a calibration/ranking-and-retention condition, none of these quantities can be substituted for another.

3.7 Presence, novelty, and reactivation

The events D_k^{gen} and D_k^{acq} test *presence*; the parent-inclusive union can satisfy them using an old instance without creating a new one. Let $\mathcal{V}_{k-1}^{\mathrm{gen}} = \bigcup_{t<k} C_t$ and $\mathcal{V}_{k-1}^{\mathrm{acq}} = \bigcup_{t<k} T_t$ denote previously generated and acquired phenotype sets, interpreted under the declared hash/phenotype compatibility rule. For a mathematically new, valid offspring of current true hardness at least κ, exclude both previously generated phenotypes and the *current* parent population (including the initial roots of cycle 0):

$$D_k^{\mathrm{new,gen}} = \left\{\exists x \in O_k \setminus (\mathcal{V}_{k-1}^{\mathrm{gen}} \cup P_k) : h_k^{\mathrm{BQ}}(x) \geq \kappa\right\}.$$

This exclusion is essential when a valid mutation returns the identical parent phenotype or when its child hash duplicates another current parent. It does not alter the parent-inclusive presence event D_k^{gen}. A new acquisition under the cumulative acquisition ledger is

$$D_k^{\mathrm{new,acq}} = \{\exists x \in T_k \setminus \mathcal{V}_{k-1}^{\mathrm{acq}} : h_k^{\mathrm{BQ}}(x) \geq \kappa\}.$$

A newly generated *and* acquired hard instance satisfies

$$D_k^{\text{new,gen}\cap\text{acq}} = \left\{\exists x \in (O_k \cap T_k) \setminus (\mathcal{V}_{k-1}^{\text{gen}} \cup P_k) : h_k^{\text{BQ}}(x) \geq \kappa\right\}.$$

This is a stricter event than $D_k^{\text{new,acq}}$: an old parent may be selected for the first time, but is not thereby a newly generated phenotype. An old instance may become hard under q_k after being non-hard under an earlier checkpoint; this is *hardness reactivation*, not instance novelty. A previously unvisited hard niche is a different event again. The baseline bridge-chain theorem below concerns *presence*, not these stronger novelty events; it may not be cited as a novelty-discovery bound without additional exclusion conditions.

3.8 First hit and computational normalization

For the abstract survival lemma, index declared comparison opportunities by j; j is not automatically a BQ cycle. An opportunity may be one serialized mutation attempt, one full generation event D_k^{gen}, or one acquisition event D_k^{acq}, but the choice must be fixed before the comparison. Let

$$\tau = \inf\{j \geq 1 : Y_j = 1\}, \qquad C_\tau = \sum_{j=1}^{\tau} c_j,$$

where Y_j is the declared true-hardness discovery event and $c_j > 0$ is the normalized compute cost through resolving opportunity j. The budget-B discovery event is $\{C_\tau \leq B\}$. The cost must include the work that differs between methods (structural generation, HGB40 inference, archive operations, and BQ evaluation/verification, as applicable). Counting generations alone is not computational normalization. The elementary bound below uses deterministic opportunity costs c_i; random or method-dependent costs require a cost-aware hazard analysis and are not silently replaced by generation counts.

For a policy π, define its conditional hit hazard

$$q_j^\pi = \mathbb{P}(Y_j = 1 \mid \mathcal{F}_{j-1}^\pi),$$

where $\mathcal{F}_{j-1}^\pi$ is the complete pre-opportunity history. On a no-hit history, a lower bound on q_j^π is a reachability certificate. A Markov state representation is one way to calculate this hazard, but the history-based definition does not require Markovianity.

3.9 Abstraction versus verified BLADE-NCO context

Supplied component	Representation here	Unspecified for a BLADE-specific proof
FOGA-11 mutation	Eleven extended kernels $\widetilde{K}_{k,o}$ on $\overline{\mathcal{X}}$	Exact kernel laws over stochastic parameters, rejection, and support bounds
Inherited parents; Elite-QD1000 (P=T)	P_k, measurable partition g, deterministic HGB/QD selection, $P_{k+1} = T_k$	Distributional survival/reachability bounds under a declared ensemble
HGB40 and BQ co-adaptation	Prediction $\widehat{h}_k$, true BQ hardness h_k^{BQ}, solver state q_k	Scientific threshold, calibration/retention bridge, and matched validation cost
Fixed generator and HAC	Comparator proposal kernels	Exact proposal laws, gradient access, stopping rules, and per-opportunity compute
Finite budget	Cumulative cost C_τ	Which compute resources are charged and how variable costs are normalized

Accordingly, the formal model is a deliberately explicit abstraction. The verified schedule and selection mapping do not assert unobserved transition probabilities. The remaining implementation-to-theory gaps are recorded as open assumptions in `assumption_register.md`.

3.10 Niche aggregation and the Markov condition

Suppose a Markov kernel Q on a full state space is aggregated by a finite measurable partition $\{B_1, \ldots, B_L\}$. The strong-lumpability equalities

$$Q(s, B_j) = Q(s', B_j) \quad \text{for every } s, s' \in B_i \text{ and every } i, j$$

are sufficient for the aggregate block process to be Markov for every initial law. They are not necessary for a particular initial law. A weaker, distribution-dependent Markov aggregation may hold when the conditional within-block law induced by that initial distribution and history makes the averaged transition to each B_j history-independent (possibly with a time-dependent aggregate kernel). Such a claim must name that law and prove the conditional equality; it is not obtained merely by averaging once.

The counterexample in Section 6.1 shows failure of the strong condition and therefore invalidates an implementation-independent Markov claim based on niche alone. It does not rule out a special weak aggregation for a specified initial distribution. In S6-AE, niche summaries hide parent identity, within-niche HGB order, candidate age, solver/archive state, and the deterministic hash tie-break, so strong lumpability is not established and a distribution-dependent alternative remains OPEN.

4 Candidate propositions and Stage 2 results

4.1 Stage 2 Proposition: deterministic niche persistence

[Mathematical status: PROVED] [Implementation applicability: VERIFIED for S6-AE] [Scientific novelty: NOT CLAIMED]

Proposition 1 (Occupied-niche persistence under the verified selector)**.** *Fix a niche map g. At cycle k, suppose (i) every member of P_k is in C_k; (ii) C_k is finite and unique by instance hash, with equal-hash candidates niche-compatible under the declared union semantics; (iii) $|\mathsf{S}(C_k)| \le m \le |C_k|$; (iv) $T_k = T_m(C_k)$ contains the complete per-niche floor $Q(C_k)$; and (v) $P_{k+1} = T_k$. Then*

$$\mathsf{S}(P_{k+1}) = \mathsf{S}(C_k) \supseteq \mathsf{S}(P_k).$$

If $C_k = \mathrm{unique}(P_k \cup O_k)$ under the same fixed map, then more precisely

$$\mathsf{S}(P_{k+1}) = \mathsf{S}(P_k) \cup \mathsf{S}(O_k).$$

The statement remains true when some or all scheduled mutation attempts are rejected, because rejection contributes no child and causes no resampling.

The assumptions match the inspected S6-AE selection path: the corrected C00 union and later parent–child unions contain all old parents; QD4, normalization, centroids, and tie rules are fixed; at most 1,000 niches are occupied while $m = 10{,}000$; selection contains exactly one current-HGB elite from every occupied candidate niche; and $P = T$. The completed run's strict checks verify these operational invariants. This proves a niche-set invariant, not parent-hash survival, within-niche structural equivalence, or future reachability.

Lemma 1 (Exact membership test for a candidate)**.** *Under the deterministic selector in the formal model, a candidate survives if and only if it is its niche's score elite (including the declared hash tie rule), or it lies within the first $m - |Q(C)|$ positions of the global order on $C \setminus Q(C)$. Hence neither membership in P_k nor occupancy of its niche is sufficient for that candidate to survive.*

4.2 Stage 2 theorem candidate: a certified bridge chain

[Mathematical status: PROVED under explicit premises] [Implementation applicability: PARTIAL/OPEN] [Scientific novelty: NOT CLAIMED]

Proposition 2 (Finite-budget causal bridge-chain bound)**.** *Fix a starting cycle k and an integer $L \ge 1$. Let the bridge sets $B_0, \dots, B_L$ be selected by the predictable, non-anticipating rule in the formal model. Suppose E_0 holds almost surely. Let $H_0 = E_0$ and $H_{\ell+1} = H_\ell \cap G_\ell \cap R_\ell$. For every $\ell = 0, \dots, L-1$ and every positive-probability pre-generation history in H_ℓ, assume*

$$\mathbb{P}(G_\ell \mid \mathcal{F}^{\mathrm{pre}}_{k+\ell}) \ge a_\ell.$$

On the same histories for which the conditional generation probability is positive, assume the pre-generation-conditioned *retention bound*

$$\mathbb{P}(R_\ell \mid \mathcal{F}^{\mathrm{pre}}_{k+\ell}, G_\ell) \ge \rho_\ell.$$

Here the second conditional probability averages over the still-unresolved mutation batch and any other explicitly modeled random inputs; it does not condition on the fully realized scored candidate pool. Alternatively, one may assume the stronger deterministic certificate: for every admissible generated union containing a bridge member, the selector retains a member of the designated bridge set, giving $\rho_\ell = 1$.

Assume the final, causally chosen bridge set satisfies $B_L \subseteq \mathcal{H}_{k+L-1}$ *almost surely for the entering checkpoint* q_{k+L-1}*, without looking ahead to future solver states or the final random mutation batch. Define the exact certified-chain success events*

$$D^{\mathrm{chain,gen}} = H_{L-1} \cap G_{L-1}, \qquad D^{\mathrm{chain,acq}} = H_{L-1} \cap G_{L-1} \cap R_{L-1}.$$

These are subsets of the ordinary final-cycle presence events D^{gen}_{k+L-1} *and* D^{acq}_{k+L-1} *respectively. Then*

$$\mathbb{P}(D^{\mathrm{chain,gen}}) \geq \prod_{\ell=0}^{L-1} a_\ell \prod_{\ell=0}^{L-2} \rho_\ell,$$
$$\mathbb{P}(D^{\mathrm{chain,acq}}) \geq \prod_{\ell=0}^{L-1} a_\ell \rho_\ell.$$

No independence is required. If the declared deterministic charged costs $c_k, \ldots, c_{k+L-1}$ *sum to at most budget* B*, these are finite-budget* presence *bounds, not new-instance discovery bounds.*

One sufficient operator-specific source for a_ℓ *is a pre-generation-measurable choice of a retained parent* $x \in B_\ell$ *and a designated operator* $o_\ell \in \mathcal{O}$ *such that*

$$\widetilde{K}_{k+\ell,o_\ell}(x, B_{\ell+1} \mid S_{k+\ell}) \geq a_\ell$$

on every admissible history, where the conditional kernel is defined under the declared non-frozen random ensemble and includes rejection at ∂*. The all-operator deterministic batch attempts the designated operator without any random parent/operator-selection factor. To infer* G_ℓ *from the accepted operator output, equal hashes must identify the same phenotype on that pool or at least preserve membership in* $B_{\ell+1}$*; niche-label compatibility alone does not suffice for the bridge theorem.*

The deterministic retention certificate is directly testable using the actual HGB/QD rank rule. Specifically, $\rho_\ell = 1$ follows if every admissible realized scored candidate union containing $B_{\ell+1}$ retains a bridge member either as the tie-broken per-niche score elite or through global fill (Lemma 1). After the complete scored union is known, the retention probability is the indicator $\mathbf{1}_{R_\ell}$, not an unexplained fractional probability. A fractional $0 < \rho_\ell < 1$ requires a clearly declared ensemble and a lower bound on retention conditioned only on pre-batch information and the event G_ℓ. Within-niche transition equivalence of possible elites would be one *additional* way to certify future reachability, not a property guaranteed by QD occupancy or hardness prediction accuracy.

For S6-AE, the all-operator schedule, rejection handling, capacity inequality, selector form, and $P = T$ map are verified. Positive, history-uniform a_ℓ values, scientifically meaningful *causally chosen* bridge sets, and deterministic or probabilistic ρ_ℓ bounds are OPEN. Changing BQ checkpoints are allowed, but terminal set inclusion must be verified against the checkpoint already available before the final batch. Freezing every bridge set at cycle k would require the stronger uniform terminal-inclusion guarantee over the possible future solver trajectories.

Repeated certified windows. [**Mathematical status: PROVED under regeneration**] [**Implementation applicability: OPEN**] [**Scientific novelty: NOT CLAIMED**] If after every no-hit L-cycle block the next block again starts with E_0 and satisfies the same conditional premises, set $p = \prod a_\ell \prod \rho_\ell$ for the desired discovery level and let C bound a block's charged compute. The standard survival argument gives $\mathbb{P}(C_\tau \le rC) \ge 1-(1-p)^r$ and $\mathbb{E}[C_\tau] \le C/p$. The regeneration premise is not verified for S6-AE; without it, the one-block finite-budget bound is the strongest conclusion.

4.3 Mechanism-level comparator boundary

[**Mathematical status: CONDITIONAL FRAMEWORK**] [**Implementation applicability: OPEN**] [**Scientific novelty: OPEN**] The bridge theorem permits a fair structural comparison only after each method receives its own transition and cost model.

- A reset evolution strategy can erase intermediate bridge sets between rounds, so a long path certificate may fail, but a reset distribution can still have larger direct mass on the true-hard region.
- A stationary generator has no inherited structural chain; its discovery law is determined by its sampling mass on the cycle-specific hard region and its per-sample cost. It can beat BLADE when that mass is sufficiently large.
- A HAC-style method with random restarts requires a kernel on its full gradient state, a restart law, stopping rules, and charged forward/backward costs. One gradient step is not equated with one FOGA attempt or one full 110,000-attempt generation. The Stage 1 path counterexample shows only that a useful gradient can win in a toy model; it is not a HAC superiority theorem.

Proposition 1 is a structural-coverage statement. Proposition 2 concerns true-hardness generation or acquisition only when its final bridge set lies in the declared true BQ hard region and, for acquisition, survives selection. Neither QD occupancy nor HGB score alone supplies that implication.

4.4 Candidate 1: finite-window reachability comparison

[**Mathematical status: CONDITIONAL**] [**Implementation applicability: PARTIAL**] [**Scientific novelty: NOT CLAIMED**] The following is the strongest surviving candidate.

Proposition 3 (Memory advantage on a stable target window)**.** *Suppose the true hard region attached to the declared discovery event is constant, $\mathcal{H}_j = \mathcal{H}$, for a window of W comparison opportunities and each opportunity has the same deterministic compute cost $c > 0$. On every no-hit history in that window, suppose a memory policy M has conditional hit hazard at least p_M, and a comparator G has conditional hit hazard at most p_G, with $p_M > p_G$. Then, for $1 \le n \le W$,*

$$\mathbb{P}(\tau_M \le n) \ge 1-(1-p_M)^n > 1-(1-p_G)^n \ge \mathbb{P}(\tau_G \le n),$$

and, for $2 \le n \le W$,

$$\mathbb{E}[c\min(\tau_M, n)] \le c\sum_{j=0}^{n-1}(1-p_M)^j < c\sum_{j=0}^{n-1}(1-p_G)^j \le \mathbb{E}[c\min(\tau_G, n)].$$

For a generic policy that actually samples a parent, one sufficient derivation of p_M is the following: on each no-hit history, with conditional probability at least r a useful retained parent is available;

given that event it is selected with probability at least γ; conditional on that selection, its mutation/evaluation reaches the true current hard region with probability at least α. Then $p_M \geq r\gamma\alpha$. No independence among these events is assumed; the bounds are conditional in the stated order.

That random-selection factorization is not the S6-AE schedule. For the verified deterministic Cartesian batch, every retained parent receives all 11 operators. If r lower-bounds availability of a useful parent and α_{batch} lower-bounds, under the explicitly declared RNG ensemble, the conditional probability that its scheduled operator batch generates the declared hit, then the appropriate sufficient bound is $p_M \geq r\alpha_{\text{batch}}$. With roots and all seeds conditioned to their frozen values, this is a statement about a realized batch, not an estimated random parent/operator-selection probability.

Importance. It turns a structural-memory claim into a finite-budget, equal-cost comparison that explicitly accounts for parent access and mutation reachability. It also makes clear that survival alone is not a sufficient mechanism.

Assumptions and realism. A piecewise-stable target window and equal-cost opportunities are useful first approximations. In S6-AE the BQ checkpoint is fixed during mutation/acquisition in cycle k and changes only after T_k is frozen and labeled, but comparison windows spanning adaptations can still move $\mathcal{H}_k$. HGB40, archive, exact-label, and BQ-training costs also differ from a single mutation attempt. The deterministic 10,000-by-11 schedule and $P = T$ update are verified; the history-uniform $r\alpha_{\text{batch}}$ bound, scientific threshold, and matched comparator cost/hazard remain [**Implementation applicability: OPEN**].

Proof approach. Induction on the no-hit survival probability, then the identity expressing expected truncated hitting time as the sum of survival probabilities. It requires no conditional independence and no Markov approximation. Proof details are in `proof_attempts.tex`.

Counterexample and HAC distinction. If memory's conditional hazard is lower because its retained parents are stale, or if a gradient proposal has larger hazard per compute unit, the premise fails. The exact examples in `counterexamples.tex` show both. Thus this is a conditional criterion that can distinguish BLADE from fixed generation or HAC only after their hazards and costs are measured; it does not establish that BLADE wins.

4.5 Candidate 2: inherited memory never hurts

[**Mathematical status: DISPROVED**] [**Implementation applicability: NOT APPLICABLE**] [**Scientific novelty: NOT CLAIMED**] The proposed blanket claim is that retaining historical parents cannot increase first-hitting cost relative to a fixed generator. This is false even with two structural states and one target change. The stale-memory example in Section 6.1 has expected cost 100 for inherited memory versus 2 for a fixed uniform generator. The failure is not a technical gap in Candidate 1: the required hazard dominance is reversed.

Importance and proof approach. A universal no-harm theorem would have supported the research objective directly, but the exact counterexample rules it out. No proof is possible without

additional target-drift, forgetting, or reachability assumptions. These extra assumptions would need to be stated and independently justified; they cannot be inferred from the word "memory."

Assumptions and realism. The universal claim would require memory never to reduce true-hardness reachability per unit cost for any target shift or parent-selection history. That is not guaranteed by inherited populations and is contradicted by the stale-memory example; realism for BLADE-NCO is low absent explicit target-aware forgetting and verified hazard bounds.

Strongest proof approach. An exact finite-state counterexample proves the universal claim false. A positive restricted result would need additional drift and reachability premises, which are instead stated explicitly in Candidate 1.

HAC distinction. Nothing in this false claim compares BLADE to HAC. HAC can outperform a poorly aligned retained population, as the path example in Section 6.1 demonstrates.

4.6 Candidate 3: diversity always improves discovery

[**Mathematical status: DISPROVED**] [**Implementation applicability: NOT APPLICABLE**] [**Scientific novelty: NOT CLAIMED**] The proposed blanket claim is that preserving a more diverse population always reduces first-hitting time. In a three-state model, a diverse parent selector spends half of each equal-cost proposal on an irrelevant niche and has expected hitting cost 2, while a focused selector hits in cost 1. The exact construction is in Section 6.1.

Importance and assumptions. Diversity can preserve useful future routes, but its benefit depends on niche relevance, selection allocation, target drift, and cost. A monotonicity proof from diversity cardinality alone cannot work. This counterclaim concerns a universal guarantee, not a claim that Elite-QD1000 behaves like the toy selector.

Realism and strongest proof approach. No selection-allocation or niche-relevance condition is supplied for Elite-QD1000, so the premise is not verified for BLADE-NCO. The exact three-state construction is a complete disproof of the universal claim; a positive diversity theorem would require explicit target-access conditions. This does not rank BLADE against HAC.

HAC distinction. It provides no general ranking against HAC. With equal compute, a focused or gradient-guided policy may spend more proposals on the current reachable hard region; the actual ranking depends on transition probabilities and costs, neither of which is specified here.

5 Proof attempt for the strongest candidate

5.1 History-based survival lemma

[Mathematical status: PROVED] [Implementation applicability: PARTIAL] [Scientific novelty: NOT CLAIMED] The lemma below is a standard finite-horizon hitting bound, included as the proof engine rather than as a novelty claim.

Let Y_j be the indicator of the declared discovery event at comparison opportunity j and let $\tau = \inf\{j \geq 1 : Y_j = 1\}$. The event $\{\tau \geq j\}$ is determined by the history before opportunity j. Assume there are deterministic numbers $p_j \in [0,1]$ such that, on every pre-opportunity history with no previous hit,

$$\mathbb{P}(Y_j = 1 \mid \mathcal{F}_{j-1}) \geq p_j.$$

This is a history-uniform conditional hazard condition. It allows the target attached to the event to vary with j. The index j is deliberately not the BQ cycle index k: it may represent a serialized mutation attempt, a full generation event D_k^{gen}, or an acquisition event D_k^{acq}, provided the choice and filtration are fixed.

Lemma 2 (Non-stationary survival and cost bound)**.** *For every $n \geq 1$,*

$$\mathbb{P}(\tau > n) \leq \prod_{i=1}^{n}(1 - p_i).$$

If opportunity costs c_i are deterministic, then

$$\mathbb{E}[C_\tau] \leq \sum_{i \geq 1} c_i \prod_{j=1}^{i-1}(1 - p_j),$$

with the inequality valid also when the right side is infinite. For a budget B, assume $\sum_{i=1}^{\infty} c_i = \infty$ and let $n(B) = \max\{n \geq 0 : \sum_{i=1}^{n} c_i \leq B\}$, with empty sum zero.

$$\mathbb{P}(C_\tau \leq B) \geq 1 - \prod_{i=1}^{n(B)}(1 - p_i).$$

Proof. Write $S_i = \{\tau \geq i\}$ for survival to the start of opportunity i. Since $S_i \in \mathcal{F}_{i-1}$,

$$\begin{aligned}\mathbb{P}(S_{i+1}) &= \mathbb{E}[\mathbf{1}_{S_i}\mathbb{P}(Y_i = 0 \mid \mathcal{F}_{i-1})] \\ &\leq (1 - p_i)\mathbb{P}(S_i).\end{aligned}$$

Starting from $\mathbb{P}(S_1) = 1$ and iterating gives $\mathbb{P}(S_{n+1}) \leq \prod_{i=1}^{n}(1 - p_i)$, which is the first claim. For the cost, use the nonnegative-sum identity

$$C_\tau = \sum_{i \geq 1} c_i \mathbf{1}_{\{\tau \geq i\}}, \qquad \mathbb{E}[C_\tau] = \sum_{i \geq 1} c_i \mathbb{P}(S_i),$$

and apply the survival bound termwise. Finally, completion of any of the first $n(B)$ opportunities costs at most B, so budget success is at least the probability of a hit among those opportunities. The first claim supplies the stated lower bound. □

5.2 Deriving the window comparison

For Candidate 1, apply the survival lemma on the W-opportunity stable-target window. The memory policy satisfies $p_j = p_M$ and the comparator satisfies an upper hazard bound $q_j^G \le p_G$ on every no-hit history. The same induction, with the inequality reversed for G, yields

$$\mathbb{P}(\tau_M > j) \le (1 - p_M)^j, \qquad \mathbb{P}(\tau_G > j) \ge (1 - p_G)^j.$$

Therefore, for $n \le W$,

$$\mathbb{P}(\tau_M \le n) \ge 1 - (1 - p_M)^n > 1 - (1 - p_G)^n \ge \mathbb{P}(\tau_G \le n).$$

For $2 \le n \le W$, using $\mathbb{E}[\min(\tau, n)] = \sum_{j=0}^{n-1} \mathbb{P}(\tau > j)$ and multiplying by the common opportunity cost c gives the strict truncated compute-cost comparison in Proposition 3. At $n = 1$ both policies pay the first opportunity cost in the truncated-cost convention, so only the success-probability comparison is strict there.

5.3 Stage 2 proof of deterministic niche persistence

Proof of Proposition 1. For every $j \in \mathsf{S}(C_k)$, the intersection $C_k \cap \mathcal{N}_j$ is nonempty, so $e_j(C_k)$ exists. By construction, $e_j(C_k) \in Q(C_k) \subseteq T_k = P_{k+1}$. Hence $\mathsf{S}(C_k) \subseteq \mathsf{S}(P_{k+1})$. Conversely $T_k \subseteq C_k$, so $\mathsf{S}(P_{k+1}) \subseteq \mathsf{S}(C_k)$ and equality holds. Since $P_k \subseteq C_k$, monotonicity follows. If C_k is the unique union of P_k and O_k, applying g gives $\mathsf{S}(C_k) = \mathsf{S}(P_k) \cup \mathsf{S}(O_k)$. A rejected outcome is $\partial \notin \mathcal{X}$ and is absent from O_k, so it changes none of these inclusions. No claim about which member realizes a niche is used. □

Proof of Lemma 1. The two components $Q(C)$ and $F_m(C)$ are disjoint. The first contains exactly the declared score-and-hash winner in each occupied niche. The second is, by definition, exactly the prefix of length $m - |Q(C)|$ of the global order on the remaining candidates. Their union therefore yields the stated necessary and sufficient condition. A previous parent that loses both tests is absent from $T_m(C)$ even though another member of its niche is present. □

5.4 Stage 2.1: corrected proof of the causal bridge-chain bound

Proof of Proposition 2. By non-anticipation, H_ℓ and the two bridge sets used in cycle $k + \ell$ are known at $\mathcal{F}_{k+\ell}^{\mathrm{pre}}$; in particular the terminal set is selected before its final mutation batch. Conditional on any positive-probability pre-generation history in H_ℓ, the ordered conditional-probability identity gives

$$\begin{aligned}&\mathbb{P}(G_\ell \cap R_\ell \mid \mathcal{F}_{k+\ell}^{\mathrm{pre}}) \\ &\quad = \mathbb{P}(G_\ell \mid \mathcal{F}_{k+\ell}^{\mathrm{pre}})\, \mathbb{P}(R_\ell \mid \mathcal{F}_{k+\ell}^{\mathrm{pre}}, G_\ell) \ge a_\ell \rho_\ell.\end{aligned}$$

The event-conditioned factor is only required on histories where the first factor is positive. In the degenerate case $a_\ell = 0$, the desired lower bound is zero without any retention condition. No independence or random parent/operator selection is used.

Since $H_\ell \in \mathcal{F}_{k+\ell}^{\mathrm{pre}}$, the tower property yields

$$\mathbb{P}(H_{\ell+1}) = \mathbb{E}\big[\mathbf{1}_{H_\ell} \mathbb{P}(G_\ell \cap R_\ell \mid \mathcal{F}_{k+\ell}^{\mathrm{pre}})\big] \ge a_\ell \rho_\ell \mathbb{P}(H_\ell).$$

Starting with $\mathbb{P}(H_0) = 1$ and iterating through $\ell = L - 2$ gives

$$\mathbb{P}(H_{L-1}) \geq \prod_{\ell=0}^{L-2} a_\ell \rho_\ell.$$

In the last cycle, pre-generation target correctness is $B_L \subseteq \mathcal{H}_{k+L-1}$ almost surely on admissible histories. Thus $H_{L-1} \cap G_{L-1}$ implies a current-solver true-hard instance is present in the final candidate union, and has probability at least $\mathbb{P}(H_{L-1})a_{L-1}$. Requiring R_{L-1} as well yields the acquisition-level factor ρ_{L-1}.

For the operator-specific sufficient condition, choose a bridge parent from $P_{k+\ell} \cap B_\ell$ by a fixed measurable rule (for example the least hash, assuming the hash convention distinguishes needed phenotypes). Because all eleven operators are scheduled, the designated operator is attempted on that parent. Its full, rejection-aware conditional kernel assigns mass at least a_ℓ to $B_{\ell+1}$. Provided that hash deduplication identifies equal solver-visible phenotypes on the realized pool (or at least preserves membership in $B_{\ell+1}$), an accepted bridge result remains represented in the candidate union even if it duplicates an earlier occurrence. This establishes the generation premise.

Once the complete scored candidate union is conditioned upon, R_ℓ is deterministic and $\mathbb{P}(R_\ell \mid \mathcal{F}^{\text{score}}_{k+\ell}) = \mathbf{1}_{R_\ell}$. Therefore a pathwise elite/global-fill membership certificate guarantees $\rho_\ell = 1$. A fractional ρ_ℓ can only be a conditional average over randomness *not* already revealed by $\mathcal{F}^{\text{pre}}_{k+\ell}$, and is not implied by the deterministic selector itself. The cost statement follows when the total charged cost of the completed certified chain is bounded by B as assumed. □

For repeated certified blocks, condition on each no-hit block boundary. An explicit regeneration hypothesis must restore both the initial bridge availability and the same conditional transition/retention bounds. The history-based survival lemma then gives survival at most $(1-p)^r$ after r blocks and expected cost at most C/p if the cost of each block is bounded above by C. No regeneration claim is made for S6-AE.

5.5 Dependency audit

- **Conditional independence:** not used. The proof requires history-conditional lower/upper hazard bounds; dependence across offspring or generations is allowed if those bounds remain true after every no-hit history.
- **Markovianity:** not used in the hitting proof. It would only be needed if one chooses a finite-state transition matrix to calculate the hazards. A niche aggregation is valid for that purpose only under lumpability.
- **Mutation bounds:** the quantity α must be a lower bound on reaching the true current hard region conditional on the useful parent being selected. A graph path's existence alone is insufficient; its probability and compute cost matter. The Stage 2 bridge bound instead exposes operator-specific accepted-output mass in $\widetilde{K}_{k,o}$; rejection probability remains in the denominator and no accepted-only renormalization is permitted.
- **Parent retention and scheduling:** a lower bound r must concern availability of a genuinely useful parent at the declared opportunity. Historical presence before a target shift does not establish continued usefulness afterward. The S6-AE parent/operator schedule is deterministic: conditional on P_k, each parent receives all 11 operators. Hence no random parent-selection factor γ may be inserted; a batch reachability bound α_{batch} must be derived under the declared RNG ensemble.

- **Causal target choice and retention information:** the bridge sets must be fixed before their respective generation batches. The last target is the true hard region of the *entering* solver checkpoint. Selection is deterministic after scoring; a nontrivial fractional ρ_ℓ must be conditioned on a coarser, pre-generation sigma-algebra, or replaced with a pathwise rank certificate.
- **Hash-quotient semantics:** niche occupancy only needs equal-hash candidates to share a niche, but bridge-chain propagation additionally needs equal hashes to identify equivalent phenotypes or at least to preserve membership in every chosen bridge set.
- **Surrogate accuracy:** no accuracy is needed for the abstract survival lemma. To pass from D_k^{gen} to D_k^{acq}, a retention condition connecting $\widehat{h}_k$ and the QD rule to true h_k^{BQ} and future reachability is essential. Perfect current-hardness prediction alone is insufficient when tied candidates have different mutation futures. Without such a bridge, the implementation-level step is OPEN.
- **Capacity and selection:** niche persistence uses $|\mathsf{S}(C_k)| \le m$ and inclusion of the complete floor $Q(C_k)$. It does not imply bridge persistence. The latter requires the exact elite/global-fill membership condition or another proved retention certificate.
- **Changing hardness:** the bridge theorem allows cycle-indexed kernels and sets, but $B_L \subseteq \mathcal{H}_{k+L-1}$ must refer to true hardness under the entering checkpoint for the final acquisition cycle. Current HGB quality cannot silently replace that condition.
- **Cost normalization:** the comparison assumes equal deterministic cost c per opportunity. If BLADE archive updates, HGB40 inference, or BQ validation change cost, opportunity-count dominance does not imply compute-cost dominance.

5.6 Failure boundary

[**Mathematical status: OPEN for a BLADE-specific dominance claim**] [**Implementation applicability: PARTIAL**] [**Scientific novelty: OPEN**] The proof does not establish that BLADE's conditional hazard is above the fixed-generator or HAC hazard. It fails to produce a BLADE result at the exact step where one would need verified bounds on r and α_{batch} (or a direct generation/acquisition hazard bound) and a comparator upper bound p_G, all in the same compute units. The counterexamples show these inequalities can reverse. No unproved transition from a structural-memory narrative to hazard dominance is made.

6 Adversarial finite-state checks

The examples below are exact toy models. They test universal statements and assumptions; none is asserted to be the BLADE-NCO implementation. Their tags report mathematical status; implementation applicability is [**Implementation applicability: NOT APPLICABLE as a universal guarantee**], and novelty is [**Scientific novelty: NOT CLAIMED**].

6.1 Stale memory after a target shift

[**Mathematical status: DISPROVED**] for the claim that historical memory cannot hurt. Take $\mathcal{X} = \{a, b\}$. Before a change the hard region is $\mathcal{H}_0 = \{a\}$; just before the first measured proposal it shifts to $\mathcal{H}_i = \{b\}$ for all $i \ge 1$. The inherited population contains only a. The memory policy mutates from a to a with probability 0.99 and to b with probability 0.01, retaining a after a

miss. Each proposal costs one unit. Its hitting time is geometric with success probability 0.01, so $\mathbb{E}[\tau_M] = 100$ and

$$\mathbb{P}(\tau_M \leq B) = 1 - 0.99^B.$$

A fixed generator that independently samples a or b uniformly every proposal has $\mathbb{E}[\tau_G] = 2$ and success probability $1 - 0.5^B$. Thus the memory policy is strictly worse at every positive integer budget. The target has undergone a single non-stationary shift; the failure is stale structural support, not an independence artifact.

6.2 Gradient adaptation can beat structural mutation

[**Mathematical status: DISPROVED**] for any universal claim that structural mutation outperforms HAC-style gradient adaptation. Let $\mathcal{X} = \{0, 1, \ldots, L\}$ and $\mathcal{H} = \{L\}$. A structural mutation from $i < L$ advances to $i+1$ with probability ε and otherwise stays at i. Starting at 0, its expected number of unit-cost proposals to hit L is L/ε. A gradient-informed comparator that advances deterministically one step per unit-cost proposal hits in L proposals. For any $0 < \varepsilon < 1$, the structural policy is slower. This is a toy landscape with an informative gradient, not a model of the actual HAC implementation; it proves that operator labels alone imply no ranking.

6.3 Diversity can allocate effort away from the target

[**Mathematical status: DISPROVED**] for the claim that increased diversity always improves first hitting time. Let $\mathcal{X} = \{s, u, g\}$, target $\mathcal{H} = \{g\}$, and one proposal cost one unit. The diverse parent set contains distinct niches s and u; parent selection is uniform. Mutation from s reaches g surely, whereas mutation from u stays at u. Hence the per-proposal hit probability is $1/2$ and expected hitting cost is 2. A focused policy that selects s every time hits in cost 1. Preserving the irrelevant niche does not help this target and reduces access to the useful parent.

6.4 Perfect current-hardness prediction need not preserve a bridge

[**Mathematical status: DISPROVED**] for the claim that perfect prediction of current BQ hardness makes score-elite retention preserve future mutation reachability. Consider states $\{a, b, g\}$ and current target $\mathcal{H} = \{g\}$; neither a nor b is itself a hit, and their current true hardness is equal: $h(a) = h(b) < \kappa$. From a, the next scheduled mutation reaches g surely; from b, mutation stays at b surely. Let the surrogate be perfect on current hardness, so $\widehat{h}(a) = h(a) = h(b) = \widehat{h}(b)$. A one-slot score-elite update must still break the tie. If its deterministic tie rule retains b, it discards a and makes g unreachable. Thus current-hardness accuracy, even perfection, does not identify equal-hardness candidates with different mutation futures. A reachability-aware rule could retain a, but perfect current-hardness prediction alone need not do so. This is not a claim that the observed S6-AE hash tie-break selected such a pair; it is an exact obstruction to the proposed bridge.

6.5 Monotone niche occupancy with progressive bridge loss

[**Mathematical status: PROVED counterexample**] to the claim that monotone niche occupancy preserves future hard-region reachability. Fix population capacity m and an integer

$2 \le L \le \min\{K-1, m-1\}$. Use occupied niches $0, 1, \ldots, L$. Initially P_0 contains a bridge a_i in niche i for each $1 \le i \le L$, together with $m - L$ distinct fillers $f_1, \ldots, f_{m-L}$ in niche 0.

At cycle $i-1$, one designated operator applied to a_i produces a valid decoy d_i in the same niche i with probability one. Every other scheduled attempt returns the rejection symbol ∂; no attempt is resampled. Set the current true hardness and a perfect current predictor so that

$$h(a_i) = \widehat{h}(a_i) = 1, \qquad h(d_i) = \widehat{h}(d_i) = 2, \qquad h(f_j) = \widehat{h}(f_j) = 3,$$

with deterministic distinct hash order where needed. Choose the current hard threshold above 3, so none of these candidates is a current hit.

In niche i, d_i replaces a_i as the score elite. Niche 0 contributes one filler to the elite floor, and the remaining $m - L - 1$ fillers occupy all global-fill slots. Thus a_i fails both tests in Lemma 1 and is discarded. All as-yet-unprocessed a_j remain the sole candidates in their niches and therefore survive. Inductively, after cycle t the selected population still occupies exactly $\{0, 1, \ldots, L\}$, while the number of retained bridges has fallen from L to $L - t$.

After cycle $L-1$, let the future true-hard region be $\{g\}$. A future designated operator reaches g surely from any a_i and cannot reach it from any d_i or f_j. The selected population now contains no a_i, so its future hit probability is zero, even though its niche occupancy never decreased. A population retaining any discarded a_i would hit surely. The failure is not surrogate error: the predictor is perfect for current hardness. It is a mismatch between current-hardness quality and future-reachability value.

This construction includes the complete elite floor and global score fill, not merely a one-slot archive. Taking, for example, $m = 10{,}000$ and $L \le 999$ respects the S6-AE capacity ratio. It is a formal counterexample, not evidence that these states or score relations occurred in S6-AE.

6.6 Niche aggregation can destroy the Markov property

[**Mathematical status: DISPROVED**] for the assertion that a niche label alone is a Markov state. Let full states x_1, x_2 both map to niche A, while z, w map to distinct niches Z, W. From x_1, the next state is z with probability 1; from x_2, the next state is w with probability 1. Both full states have the same current aggregate A, but the next aggregate is respectively Z or W. The transition probability out of aggregate A is therefore not determined by the aggregate alone (strong lumpability fails). Therefore niche aggregation is not Markov uniformly over initial laws. For a specially chosen initial distribution, an averaged transition might nevertheless be Markov; that weaker distribution-dependent claim would require the conditional within-A mixture to remain history-independent. The example does not disprove every such weak aggregation. A proof using only niche counts is invalid unless it proves the relevant strong or distribution-dependent condition, or augments the state with the hidden structural/solver information.

6.7 Hard-instance presence is not novel discovery

[**Mathematical status: PROVED definitional counterexample**] to treating the parent-inclusive generation hit as new-instance discovery. Let $P_k = \{x\}$, $h_k^{\mathrm{BQ}}(x) \ge \kappa$, and all scheduled child attempts be rejected, so $O_k = \varnothing$ and $C_k = \{x\}$. Then D_k^{gen} holds, while $D_k^{\mathrm{new,gen}}$ is false whenever x was previously generated. For a separate cycle-0 edge case, suppose $P_0 = \{x\}$ and a

valid, accepted mutation returns exactly x, so $O_0 = \{x\}$ before parent–child union deduplication. Although $\mathcal{V}_{-1}^{\text{gen}} = \varnothing$, this offspring is not a novel phenotype because $x \in P_0$. The corrected event $D_0^{\text{new,gen}}$ therefore excludes it. After every structural niche is already occupied, a later offspring can still be a new, distinct hard phenotype within an old niche. Therefore newly occupied niches, new hard phenotypes, acquired novelty, and hardness reactivation must not be conflated.

7 Stage 3A: solver-relative hardness tracking

This section studies whether the verified HGB40 plus Elite-QD1000 ($P = T$) mechanism tracks high current BQ hardness inside each fixed niche as the solver changes. It does not identify current hardness with future reachability or downstream training value, and it does not compare BLADE to HAC. All maxima below are over finite realized candidate sets; no finiteness of the ambient instance space is assumed.

7.1 Tracking objects and information order

Let M_k be the BQ solver entering cycle k and define, on the same scale as the fitted HGB target,

$$f_k(x) = \log(1 + L_{\text{BQ}}(M_k, x)).$$

Let $\widehat{f}_k$ be the median HGB40 prediction. For every $j \in \mathsf{S}(C_k)$ write

$$E_{k,j} = C_k \cap \mathcal{N}_j, \qquad e_{k,j} = \operatorname*{arg\,max}^{\text{score/hash}}_{x \in E_{k,j}} \widehat{f}_k(x),$$

where "score/hash" means descending prediction followed by ascending instance hash. Thus $e_{k,j}$ is a single, deterministic member of the per-niche floor Q_k, and

$$e_{k,j} \in Q_k \subseteq T_k = P_{k+1}.$$

The prediction and selection use M_k. Exact labels on T_k and adaptation to M_{k+1} occur only after this selection. Any inequality involving $f_{k+1}(e_{k,j})$ is therefore an ex post tracking statement under a declared drift premise, not a claim that cycle k used future information.

The quantities kept distinct in this section are: current true hardness f_k, current prediction $\widehat{f}_k$, within-niche selected quality, solver-induced drift $f_{k+1} - f_k$, best hardness among all historically generated candidates, future mutation reachability, and downstream BQ training value. No implication among the last three is assumed.

7.2 One-step approximate within-niche maximization

[Mathematical status: PROVED under a finite-pool error bound] [Implementation applicability: PARTIAL] [Scientific novelty: STANDARD]

Proposition 4 (Per-niche approximate argmax)**.** *Fix k and $j \in \mathsf{S}(C_k)$. Suppose $\epsilon_{k,j} \geq 0$ and*

$$|\widehat{f}_k(x) - f_k(x)| \leq \epsilon_{k,j} \qquad \textit{for every } x \in E_{k,j}.$$

Then

$$f_k(e_{k,j}) \geq \max_{x \in E_{k,j}} f_k(x) - 2\epsilon_{k,j}.$$

If $P_k \cap \mathcal{N}_j \neq \varnothing$, parent inclusion in the candidate union also gives

$$f_k(e_{k,j}) \geq \max_{x \in P_k \cap \mathcal{N}_j} f_k(x) - 2\epsilon_{k,j}.$$

Proof. Because $E_{k,j}$ is finite and nonempty, choose $x^* \in \arg\max_{x \in E_{k,j}} f_k(x)$. The selected candidate maximizes the prediction even when the prediction maximum is tied, so

$$\begin{aligned} f_k(e_{k,j}) &\geq \widehat{f}_k(e_{k,j}) - \epsilon_{k,j} \\ &\geq \widehat{f}_k(x^*) - \epsilon_{k,j} \\ &\geq f_k(x^*) - 2\epsilon_{k,j}. \end{aligned}$$

The parent-relative result follows from $P_k \cap \mathcal{N}_j \subseteq E_{k,j}$. □

The deterministic hash rule makes the selected maximizer single-valued but does not change the proof. An empty candidate niche has no elite and no maximum, so the statement is restricted to $j \in \mathsf{S}(C_k)$. The parent-relative form additionally requires that the parent niche be nonempty. Hash deduplication must preserve the phenotype, score, and niche identity used by the statement on the realized pool. The additive 2ϵ guarantee is on the $\log(1 + \mathrm{LBQ})$ scale. On the raw scale it implies only

$$L_{\mathrm{BQ}}(M_k, e_{k,j}) \geq e^{-2\epsilon_{k,j}} \left(1 + \max_{x \in E_{k,j}} L_{\mathrm{BQ}}(M_k, x)\right) - 1,$$

not an additive raw-LBQ bound of the same size.

This is the standard approximate-argmax inequality. Elite-QD makes the chosen approximate maximizer survive into P_{k+1}; that implementation fact is useful, but it does not make the inequality itself novel. S6-AE verifies the score scale, finite candidate pool, score/hash rule, and $P = T$. It does not verify the uniform error premise on $E_{k,j}$, because exact current LBQ is not obtained for the complete candidate union.

7.3 One-step tracking across solver adaptation

[Mathematical status: PROVED under local error and drift bounds] [Implementation applicability: PARTIAL/OPEN] [Scientific novelty: STANDARD]

Proposition 5 (Drift-adjusted per-niche tracking)**.** *Under Proposition 4, suppose additionally that*

$$|f_{k+1}(x) - f_k(x)| \leq \delta_{k,j} \qquad \text{for every } x \in E_{k,j},$$

where $\delta_{k,j} \geq 0$. Then

$$f_{k+1}(e_{k,j}) \geq \max_{x \in E_{k,j}} f_{k+1}(x) - 2\epsilon_{k,j} - 2\delta_{k,j}.$$

Consequently, whenever $P_k \cap \mathcal{N}_j$ is nonempty,

$$f_{k+1}(e_{k,j}) \geq \max_{x \in P_k \cap \mathcal{N}_j} f_{k+1}(x) - 2\epsilon_{k,j} - 2\delta_{k,j}.$$

For only the parent-relative conclusion it is sufficient that the drift bound hold on $(P_k \cap \mathcal{N}_j) \cup \{e_{k,j}\}$.

Proof. The drift bound and Proposition 4 give

$$\begin{aligned} f_{k+1}(e_{k,j}) &\geq f_k(e_{k,j}) - \delta_{k,j} \\ &\geq \max_{x \in E_{k,j}} f_k(x) - 2\epsilon_{k,j} - \delta_{k,j} \\ &\geq \max_{x \in E_{k,j}} f_{k+1}(x) - 2\epsilon_{k,j} - 2\delta_{k,j}. \end{aligned}$$

Restricting the comparison set to old parents proves the second claim. The same calculation uses only the smaller stated set for the parent-relative version. □

The guarantee says that an elite selected under M_k remains competitive, after one update, with candidates that were already available in cycle k. It does not compare against new cycle-$(k+1)$ offspring, show that $e_{k,j}$ remains the next HGB elite, preserve its individual hash beyond P_{k+1}, or control future reachability or training value. The drift bound is local to an explicitly named finite set; no global bound on the measurable TSP space is needed. Even this local bound is not currently certified by S6-AE.

7.4 Historical benchmark and multicycle regret

For a finite horizon define the latent generated-candidate ledger

$$\mathcal{U}_k = \bigcup_{t=0}^{k} C_t.$$

For $j \in \mathsf{S}(\mathcal{U}_k)$ define

$$V_{k,j} = \max_{x \in \mathcal{U}_k \cap \mathcal{N}_j} f_k(x), \qquad R_{k,j} = V_{k,j} - f_k(e_{k,j}).$$

By deterministic niche persistence, once niche j has appeared by cycle k, it is occupied in P_{k+1} and hence has an elite in every later cycle. The benchmark $V_{k,j}$ is theoretical: S6-AE neither retains every member of $\mathcal{U}_k$ as a candidate nor evaluates every historical candidate under M_k.

Define the current unavailability penalty

$$A_{k,j} = \left[\max_{x \in \mathcal{U}_{k-1} \cap \mathcal{N}_j} f_k(x) - \max_{x \in C_k \cap \mathcal{N}_j} f_k(x) \right]_+,$$

with $A_{k,j} = 0$ if the historical set is empty.

Proposition 6 (Exact source of historical tracking loss)**.** *Under the prediction-error premise of Proposition 4,*

$$R_{k,j} \leq 2\epsilon_{k,j} + A_{k,j}.$$

Proof. Write $B_{k,j} = \max_{x \in C_k \cap \mathcal{N}_j} f_k(x)$. Since $\mathcal{U}_k = \mathcal{U}_{k-1} \cup C_k$, $V_{k,j} \leq B_{k,j} + A_{k,j}$. Proposition 4 gives $f_k(e_{k,j}) \geq B_{k,j} - 2\epsilon_{k,j}$, and subtraction proves the claim. □

This non-accumulating decomposition is informative only if $A_{k,j}$ can be controlled. Parent–child union makes the immediately previous elite available, which yields the following recursive control under a much stronger historical drift premise.

[Mathematical status: PROVED under full historical-set drift] [Implementation applicability: OPEN] [Scientific novelty: STANDARD]

Proposition 7 (Historical per-niche regret recurrence)**.** *Suppose niche j was occupied by cycle $k-1$. Assume:*

1. *the prediction-error premise holds on $C_k \cap \mathcal{N}_j$;*
2. *$e_{k-1,j} \in P_k \subseteq C_k$; and*
3. *for every $x \in \mathcal{U}_{k-1} \cap \mathcal{N}_j$,*

$$|f_k(x) - f_{k-1}(x)| \le \delta_{k-1,j}.$$

Then

$$R_{k,j} \le R_{k-1,j} + 2\epsilon_{k,j} + 2\delta_{k-1,j}.$$

If τ_j is the first cycle in which niche j appears, then

$$R_{\tau_j,j} \le 2\epsilon_{\tau_j,j},$$

and for every $k \ge \tau_j$,

$$R_{k,j} \le 2\sum_{s=\tau_j}^{k} \epsilon_{s,j} + 2\sum_{s=\tau_j}^{k-1} \delta_{s,j}.$$

Proof. Let $B_{k,j} = \max_{C_k \cap \mathcal{N}_j} f_k$. The retained prior elite gives $B_{k,j} \ge f_k(e_{k-1,j})$. Uniform drift over the historical set gives

$$\begin{aligned}
\max_{x \in \mathcal{U}_{k-1} \cap \mathcal{N}_j} f_k(x) &\le V_{k-1,j} + \delta_{k-1,j} \\
&= f_{k-1}(e_{k-1,j}) + R_{k-1,j} + \delta_{k-1,j} \\
&\le f_k(e_{k-1,j}) + R_{k-1,j} + 2\delta_{k-1,j} \\
&\le B_{k,j} + R_{k-1,j} + 2\delta_{k-1,j}.
\end{aligned}$$

New candidates in $\mathcal{U}_k \setminus \mathcal{U}_{k-1}$ lie in C_k, so their maximum is at most $B_{k,j}$. Hence $V_{k,j} \le B_{k,j} + R_{k-1,j} + 2\delta_{k-1,j}$. The approximate-argmax bound $f_k(e_{k,j}) \ge B_{k,j} - 2\epsilon_{k,j}$ proves the recurrence.

At first discovery there is no older candidate in niche j, so $\mathcal{U}_{\tau_j} \cap \mathcal{N}_j = C_{\tau_j} \cap \mathcal{N}_j$ and the initial bound follows directly. Iterating the recurrence gives the stated finite-horizon sum. □

The displayed recurrence proposed for Stage 3A is therefore valid, but only under drift control on the *entire historical niche ledger*, including candidates that the algorithm discarded and never relabels. Bounded per-cycle errors and drift yield a bound that grows linearly with the horizon; they do not give bounded long-run regret. A steady-state statement would need, for example, summable error and drift, a uniform bound on $A_{k,j}$, or an operational archive that repeatedly reintroduces the historical benchmark candidates. None is verified for S6-AE.

The accumulation cannot generally be removed. In a single niche with capacity one and zero solver drift, let cycle $t \ge 1$ contain the retained x_{t-1} and a new x_t, with $f(x_t) = -2t\epsilon$. Give both candidates the tied prediction $-(2t-1)\epsilon$ by errors of magnitude ϵ, and choose hashes so that the deterministic tie rule selects x_t. Starting from $f(x_0) = 0$, the historical regret after cycle t is exactly $2t\epsilon$. Thus a constant error bound alone permits linear loss even with parent–child union and no solver drift.

7.5 What happens after all 1,000 niches are occupied

Proposition 1 makes structural coverage constant once every niche is occupied. Propositions 4–7 address a different quantity: which representative is selected and how close its current solver-relative hardness is to a declared comparator. The per-niche floor continues to replace representatives after coverage saturation. Therefore coverage can be saturated while within-niche hardness quality improves, degrades after solver adaptation, or loses future stepping stones. Global HGB fill can retain additional candidates but supplies no worst-case historical or reachability certificate.

7.6 Adversarial boundary checks

The following finite constructions are exact mathematical checks, not claims about the realized S6-AE trajectory.

A. Perfect prediction with severe solver drift. Put a, b in one niche and let $P_k = C_k = \{a, b\}$. Set $f_k(a) = 1$, $f_k(b) = 0$, and $\widehat{f}_k = f_k$, so $e_{k,j} = a$. After adaptation set $f_{k+1}(a) = 0$ and $f_{k+1}(b) = H$ for arbitrary $H > 0$. The post-update gap to the old-parent maximum is H. Perfect current prediction gives no cross-checkpoint guarantee without a drift bound.

B. Small drift with inaccurate prediction. Keep the solver fixed, so $\delta = 0$. Set $f(a) = H$, $f(b) = 0$ in one niche, but $\widehat{f}(a) = 0$ and $\widehat{f}(b) = 1$. The elite is b and its current quality gap is H, which can be arbitrarily large. Small solver drift cannot replace a candidate-set prediction or ranking guarantee.

C. Full coverage with deteriorating future reachability. The progressive bridge-loss construction in Section 6.5 retains every occupied niche while successively replacing every future-useful bridge by a higher-current- hardness decoy. It applies even with perfect current-hardness prediction and the complete global-fill mechanism. Current-hardness tracking therefore does not imply future-reachability tracking.

D. Zero tracking regret with no downstream training benefit. Let every cycle have a unique current-hardness maximizer a_k in its niche and let selection be exact, so $R_{k,j} = 0$. Define the solver adaptation map on the selected training set to leave the solver unchanged, or to move it to a state with worse loss on a separate evaluation instance z. Then tracking regret is identically zero while downstream performance is unchanged or worse. A link from selected hardness to the optimization effect of a BQ update is an additional assumption, not a consequence of elite tracking.

E. A reset method can win without elite preservation. Let the inherited population contain only a, let all scheduled mutations from a return a, and let the current hard target be $\{g\}$. A reset generator that returns g with probability one hits in one equal-cost opportunity; the inherited method never hits. Historical elite preservation does not dominate a reset or stationary proposal law with better direct target mass.

7.7 Applicability and novelty audit

Statement	Mathematical validity	S6-AE applicability	Scientific novelty
Per-niche 2ϵ quality	PROVED conditional	PARTIAL	STANDARD
One-step $2\epsilon + 2\delta$ tracking	PROVED conditional	PARTIAL/OPEN	STANDARD
Historical regret decomposition	PROVED conditional	OPEN	STANDARD
Historical regret recurrence and sum	PROVED conditional	OPEN	STANDARD
Bounded long-run tracking regret	OPEN; not implied	OPEN	NOT CLAIMED
Tracking implies reachability/training value	DISPROVED without alignment	NOT APPLICABLE as a guarantee	NOT CLAIMED

Approximate maximization under uniform value error is elementary. Adding a bounded objective drift and telescoping the resulting regret is in the standard spirit of dynamic-regret and moving-optimum analyses [7, 1, 2]. Fixed-cell elite retention is also standard QD machinery [3, 4]. The only BLADE-specific ingredients used here are the exact current-HGB niche elite, parent–child union, and $P = T$, which make the previous elite available for one recursive step. A potentially novel result would need a non-vacuous, implementation-supported error/drift certificate or a hardness–reachability– training-value relation specific to FOGA-11 and BQ adaptation. This stage establishes none of those.

7.8 Stage 3A research decision

The proved inequalities are internally valid, but their decisive premises are not currently supported by S6-AE: exact current LBQ is unavailable on the complete candidate union, and the historical recurrence requires drift bounds for discarded candidates across the entire latent ledger. Counterexamples block any replacement of those premises by coverage, perfect present-time ranking alone, or low tracking regret.

C – TRACKING_ASSUMPTIONS_NOT_JUSTIFIED.

8 Stage 3B: observable guarantees and selective-label identifiability

This section freezes every Stage 3A statement and asks a different question: what follows from the selector and the labels that S6-AE actually records, without a full-pool prediction-error bound, a solver-drift bound, or an exchangeability assumption? The answer has two parts. The selector admits an exact, deterministic exchange identity. In contrast, key true-hardness quantities are not point identified from selectively acquired labels.

8.1 Objects, total order, and observability classes

Fix a cycle k and suppress k where no ambiguity results. The candidate union C is finite and hash-deduplicated. Define the strict total order $\prec$ by descending HGB40 prediction $\widehat{f}_k$ and, for equal predictions, ascending instance hash. Let $\mathrm{Top}_r(A)$ denote the first r elements of A in this order. For S6-AE,

$$m = 10{,}000, \qquad K = 1{,}000, \qquad G = \mathrm{Top}_m(C),$$

where G is the frozen, diagnostic pure-global HGB selection. The per-niche floor Q contains one $\prec$-best member of every occupied niche, so $q = |Q| \le K \le m$. The actual acquisition and next-parent population is

$$T = P_{k+1} = Q \cup \mathrm{Top}_{m-q}(C \setminus Q).$$

All set statements below concern instance identities. The stored row order of T—floor first and global fill second—is not part of those statements.

True current-checkpoint hardness remains

$$f_k(x) = \log\bigl(1 + L_{BQ}(M_k, x)\bigr).$$

Let $V_k \subseteq C_k$ be the candidates for which a verified value of f_k under the entering checkpoint M_k is available. In the inspected S6-AE artifacts this includes every member of T_k and any current-cycle calibration member whose hash also lies in C_k. A label obtained under M_t, $t \neq k$, is not thereby a member of V_k.

It is useful to separate four information classes.

1. *Deterministically observable:* candidate identities, hashes, descriptors, niches, HGB predictions and ranks, and the sets Q, T, G.
2. *Verified true-hardness:* $f_k(x)$ only for $x \in V_k$, with checkpoint identity and hash alignment verified.
3. *Latent:* any functional that depends on an unobserved $f_k(x)$ for $x \in C_k \setminus V_k$, such as a full-pool niche maximum.
4. *Hypothetical:* quantities requiring a new sampling law, error model, hardness-range restriction, future mutation law, or training-value relation.

Write $\mathcal{O}_k$ for the first two classes. A finite-pool quantity is *point identified* from $\mathcal{O}_k$ if every completion of f_k on $C_k \setminus V_k$ that agrees with $\mathcal{O}_k$ gives the same value.

8.2 Exact Elite-QD versus global-top-m exchange

[Mathematical status: PROVED]

[Implementation applicability: VERIFIED for the inspected S6-AE selector and artifacts]

[Scientific novelty: STANDARD CONSTRAINED TOP-m EXCHANGE]

Proposition 8 (Exact protected-floor exchange)**.** *Let C have a strict total order, let $G = \mathrm{Top}_m(C)$, and let $Q \subseteq C$ satisfy $q = |Q| \le m$. If*

$$T = Q \cup \mathrm{Top}_{m-q}(C \setminus Q),$$

then

$$T \setminus G = Q \setminus G$$

and

$$|T \setminus G| = |G \setminus T| \leq |Q| \leq K.$$

Proof. There are $m - |G \cap Q| \geq m - q$ elements of $G \setminus Q$. Every one of them precedes every element of $C \setminus G$. Hence the first $m - q$ elements of $C \setminus Q$ all belong to G, so the global-fill part of T contributes nothing to $T \setminus G$. Therefore $T \setminus G = Q \setminus G$. Both T and G have cardinality m, so their two directed set differences have equal cardinality. The final inequalities are immediate. □

Thus the actual population differs from pure global selection only through the floor elites that lie outside the pure top m, together with the same number of displaced global members. The result does not state that all of Q is "forced": $Q \cap G$ costs no global slot.

Define the predicted-score opportunity cost

$$\Delta_k^{\mathrm{QD}} = \sum_{x \in G_k} \widehat{f}_k(x) - \sum_{x \in T_k} \widehat{f}_k(x).$$

Proposition 9 (Exact surrogate-score cost)**.** *Under Proposition 8,*

$$\Delta_k^{\mathrm{QD}} = \sum_{x \in G_k \setminus T_k} \widehat{f}_k(x) - \sum_{x \in Q_k \setminus G_k} \widehat{f}_k(x) \geq 0.$$

It is exactly computable from the frozen HGB scores and membership records.

Proof. The common members cancel. The two remaining sets have the same size, and every member of G_k has score at least that of every candidate outside G_k. Equivalently, G_k maximizes the sum of scores among all m-subsets of C_k. Hash tie-breaking makes membership unique but permits $\Delta_k^{\mathrm{QD}} = 0$ when boundary scores tie. □

This is a deterministic surrogate-score tradeoff, not true-LBQ regret. It does not imply a change in future reachability or downstream training value. It is an elementary exchange property of a protected set plus top-m fill, not a new QD theorem.

8.3 Read-only S6-AE artifact audit

The selection identities were replayed from the ten frozen cycles without new solver calls. Let $r_k = |T_k \setminus G_k|$ and let $D_k = G_k \setminus T_k$. The label set used in the table is

$$V_k = T_k \cup (B_k \cap C_k),$$

where B_k is the current-checkpoint calibration table. "Unlab. D" is $|D_k \setminus V_k|$.

cycle	$\lvert C_k\rvert$	$\lvert Q_k\rvert$	r_k	$\lvert V_k\rvert$	unlab. D_k	Δ_k^{QD}
C00	119781	435	41	10409	41	0.574
C01	119700	831	463	18530	421	4.415
C02	119507	991	711	14306	662	9.928
C03	119487	1000	800	12899	765	14.322
C04	119531	1000	838	12231	803	22.946
C05	119547	1000	848	11767	820	27.599
C06	119667	1000	853	11521	821	41.746
C07	119661	1000	873	11358	849	75.619
C08	119700	1000	873	11143	860	70.111
C09	119664	1000	876	10977	862	83.503

The cost is the sum on the HGB log-target scale; dividing by m gives the stored mean-score difference. All ten cycles satisfy the exact set identity. Once coverage reaches 1,000 niches at C03, $|Q_k|$ remains fixed but r_k and Δ_k^{QD} continue to change. Coverage saturation therefore fixes the number of protected constraints, not their overlap with global top m or their surrogate opportunity cost.

The artifacts suffice for the deterministic columns and for selected-set truth summaries. They do not suffice for a true-hardness comparison of T_k and G_k: every cycle has unverified members of D_k. In particular, the stored `training_exact_hardness_cost.csv` gives its G row only on naturally labeled members (9,959 at C00 and 9,124 at C09), whereas its T row has all 10,000 members. That descriptive difference is not a full-set G-versus-T estimand.

8.4 Selective-label indistinguishability

[Mathematical status: PROVED IMPOSSIBILITY]

[Implementation applicability: VERIFIED INFORMATION OBSTRUCTION]

[Scientific novelty: STANDARD SELECTIVE-LABEL RESULT]

Proposition 10 (Finite-pool non-identifiability)**.** *Suppose $C \setminus V$ is nonempty. Without a restriction relating $\widehat{f}$ to f, or otherwise restricting unobserved values of f, there exist two finite worlds with identical candidate identities, descriptors, niches, HGB predictions, Q, T, G, and labels on every member of V, but different true-hardness values on $C \setminus V$.*

If an unverified candidate z shares niche j with its selected elite, the two worlds can have arbitrarily different niche maxima and selected-elite regret. With sufficiently many unverified nonselected candidates they can also give top-tail recall one and zero for the same selected set. If $(G \setminus T) \setminus V$ is nonempty, the sign and magnitude of the true hardness sum difference between G and T are not identified.

Proof. Keep the finite pool and every observable fixed. For the niche claim, set the verified elite value to one, set every other candidate in that niche to at most one in world W_0, and set $f(z) = H > 1$ in world W_H. The selector is unchanged because it uses $\widehat{f}$, but the niche regret changes from zero to $H - 1$. Since H is unrestricted, the difference is unbounded.

For top-r recall, take r selected candidates with fixed verified value one and r unverified nonselected candidates. Give the latter value zero in W_0 and value $H > 1$ in W_H. With deterministic truth

tie-breaking, the top-r set is entirely selected in W_0 and entirely unselected in W_H.

Finally let $d \in (G \setminus T) \setminus V$. Common members of G and T cancel in their true-score sums, while every member of $T \setminus G$ is verified because $T \subseteq V$. Setting $f(d) = 0$ or $f(d) = H$ leaves all observables fixed and can reverse the comparison; increasing H makes the difference arbitrarily large. Multiple unverified displaced candidates only strengthen the construction. □

This is an identifiability statement, not evidence that the realized S6-AE selection is poor. It says that the inspected observations are compatible with sharply different unseen true-hardness landscapes.

8.5 Partial identification under a known hardness range

Suppose, additionally, that a scientifically justified finite range $L \le f_k(x) \le U$ is known for every member of the finite candidate pool. Let

$$A = T \setminus G, \qquad D = G \setminus T, \qquad D_V = D \cap V, \qquad D_U = D \setminus V.$$

Because $A \subseteq T \subseteq V$, the following interval is observable:

$$\begin{aligned} &\sum_{x \in D_V} f_k(x) + |D_U| L - \sum_{x \in A} f_k(x) \\ &\quad \le \sum_{x \in G} f_k(x) - \sum_{x \in T} f_k(x) \\ &\quad \le \sum_{x \in D_V} f_k(x) + |D_U| U - \sum_{x \in A} f_k(x). \end{aligned}$$

These bounds are sharp absent further restrictions: assign every unobserved displaced value to L or to U. A known range therefore yields partial, not point, identification. Nonnegativity of LBQ supplies $L = 0$ on the log scale, but no inspected artifact supplies a finite scientific U. The resulting upper endpoint is therefore unavailable, and the read-only lower endpoints do not determine the sign in any inspected cycle.

For a threshold τ, the acquired hard count $|\{x \in T : f_k(x) \ge \tau\}|$ is exact. Full-pool recall is only bounded by allowing every latent candidate to fall below or above τ. Exact full-pool maxima similarly require a census of the relevant niche or an additional tail/range model.

8.6 What the existing labels do and do not identify

The following selected-population quantities are exactly computable without new true-LBQ calls: the current f_k distribution on T_k; the acquired count above a predeclared threshold; current hardness of each measured per-niche elite; HGB residuals on V_k or on T_k with the conditioning set stated; T_k–G_k overlap; r_k; and Δ_k^{QD}.

They characterize selected/acquired candidates only. They do not identify a full-pool HGB error distribution, a full-pool hard-tail recall, a niche maximum over C_k, historical improvement under the current checkpoint, generalization to another checkpoint, future reachability, or downstream BQ training value. The archive calibration sample is randomized within the archive, not within the

current candidate union; after intersecting it with C_k it is not a documented probability sample of $C_k \setminus T_k$. Selected-set residuals therefore cannot be extrapolated to the union without an explicit selection/sampling model.

8.7 Minimal additional information

Different targets require different additions.

1. *Exact current T versus G comparison.* Every member of $A = T \setminus G$ is already labeled. It is enough to evaluate the missing members of $D = G \setminus T$ under the same entering checkpoint. By Proposition 8, this requires at most $K = 1,000$ extra labels, and often fewer after reusing verified current labels. This is an exact paired certificate, but it has not been completed in the frozen artifacts.
2. *Finite-pool means, tail mass, or recall.* Draw unselected candidates with known, positive inclusion probabilities, preferably stratified by niche and frozen HGB-score bins. Design-based weighting can then estimate declared aggregate quantities. Deterministic selection probabilities of zero cannot support inverse-propensity extrapolation.
3. *Niche maxima or elite regret.* Verify the entire target niche, verify a declared finite contender set with a valid local error certificate, or impose and validate a tail/range model. A small random sample alone does not identify an exact maximum.
4. *Cross-cycle claims.* Re-evaluate paired identities under the same current checkpoint or posit a separately validated drift model. Old labels are not substitutes.
5. *Training value.* A causal training comparison is required; acquisition hardness alone supplies no identification. Stage 3B does not propose or run such a comparison.

A minimal future verification protocol, if separately authorized, would freeze C, Q, T, G before truth evaluation; predeclare the checkpoint, estimand, strata, and inclusion probabilities; label all still-unlabeled $G \setminus T$ for the exact paired selector comparison; add a probability sample from $C \setminus (T \cup G)$ only for full-pool aggregate estimands; preserve hash joins and nonresponse accounting; and report design-based uncertainty. This is a protocol specification, not an experiment performed here.

8.8 Novelty and Stage 3B decision

The protected-floor exchange and its K-label paired certificate are useful implementation audits, but the proof is elementary constrained top-m selection. QD already formalizes retaining high-quality representatives across behavioral regions. Selective-label non-identifiability, partial identification under bounded outcomes, and propensity-based learning under partial feedback are established themes [8, 9, 10]. Surrogate-assisted evolutionary computation likewise treats exact-evaluation management as a core issue rather than equating surrogate scores with truth [11].

The BLADE-specific contribution here is a precise audit map: because the actual T/G symmetric difference is at most K, an exact same-checkpoint comparison needs labels only for the missing displaced set, not the roughly 120,000-member union. This is practically useful but not, by itself, a publication-level theorem. Existing frozen labels do not complete that certificate and cannot identify the broader full-pool claims.

The Stage 3B research decision is

B – IDENTIFIABILITY_LIMIT_CONFIRMED.

No Stage 3C, new true-LBQ call, mutation, evolution, HGB training, BQ update, or algorithm modification is initiated.

9 Research assessment

Strongest candidate. The strongest Stage 2 candidate is Proposition 2. It turns a multicycle structural path into an auditable finite-budget bound using operator-specific accepted-transition masses and explicit selection-retention certificates. Unlike Proposition 3, it does not begin by assuming an unexplained overall hazard advantage. Proposition 1 supplies a verified structural invariant, but the progressive-loss counterexample shows why that invariant cannot supply the bridge-retention premise by itself.

Limitations and novelty. The product/conditioning argument and archive occupancy fact are standard and are not claimed as new theory. Scientific content would have to come from deriving non-vacuous FOGA-11 kernel bounds, proving that current-HGB elites retain future-relevant bridges despite BQ drift, and obtaining a strict compute-normalized separation from specified comparators. The inspected implementation verifies the deterministic operator schedule, rejection without resampling, population update, capacity relation, and $P = T$, but not those probabilistic or score–reachability premises. The result therefore does not assert a BLADE advantage.

Stage 2 verifies the S6-AE batch schedule and population update rule, but not the required probabilistic hazard bounds, a scientific hard-region threshold, or a matched HAC transition/cost law. Therefore the mathematical comparison remains conditional, its BLADE applicability remains partial, and its novelty remains open.

Testability. The structural persistence proposition is checkable directly from the selector definition and existing artifacts. The bridge theorem would require certified support or lower bounds for accepted operator transitions between declared bridge sets and proof that the selector retains those sets. The inspected artifacts record proposals, rejections, scores, and memberships, but a single frozen trajectory cannot establish history-uniform kernel bounds.

[Implementation applicability: OPEN]

No new evaluations, training, or evolutionary runs were performed for this note.

Unresolved issues. The decisive open issue is whether current-hardness selection preserves representatives with high future-reachability value often enough to complete a nontrivial bridge chain. Exact operator laws under a declared ensemble, score–bridge alignment, target sets, and matched fixed-generator/reset/HAC cost models remain necessary before a BLADE-specific comparison theorem can be claimed.

Stage 3A tracking assessment. The one-step 2ϵ quality guarantee and $2\epsilon + 2\delta$ solver-drift guarantee are mathematically valid but standard. Parent–child union and $P = T$ keep the immediately previous niche elite available, which supports a historical regret recurrence. That recurrence requires current-error bounds on the complete candidate niche and drift bounds on every candidate

in the latent historical niche ledger; S6-AE does not label those complete sets. Bounded per-cycle terms permit linear regret, not a steady-state guarantee. The Stage 3A decision is therefore **C – TRACKING_ASSUMPTIONS_NOT_JUSTIFIED**.

Stage 3B identifiability assessment. The verified selector satisfies $T_k \setminus G_k = Q_k \setminus G_k$ and exchanges at most 1,000 members with pure global HGB Top10k. Its predicted-score cost is non-negative and exactly observable. True current hardness, however, is selective: all of T_k is labeled, but most displaced members of $G_k \setminus T_k$ and most of the candidate union are not. Two truth assignments can therefore agree on every saved observable while reversing the full-pool conclusions. This is an information limit, not evidence of poor realized selection, and the Stage 3B decision is **B – IDENTIFIABILITY_LIMIT_CONFIRMED**.

10 Stopping point

Stage 1 formalization, Stage 1.5 alignment, Stage 2 persistence/reachability, Stage 2.1 independent mathematical corrections, Stage 3A solver-relative hardness tracking, and Stage 3B observability analysis are recorded. The revised bridge theorem is a conditional, pre-generation-information certificate, not an implemented or empirically verified BLADE superiority result. The Stage 3A regret recurrence is likewise conditional and does not establish bounded long-run tracking. Stage 3B specifies, but does not run, a minimal same-checkpoint verification protocol. No experiment, algorithm change, Oracle/BQ evaluation, retraining, adaptation, evolution, or Stage 3C development is included.

References

[1] D.-C. Dang, T. Jansen, and P. K. Lehre, "Populations Can Be Essential in Tracking Dynamic Optima," *Algorithmica*, vol. 78, no. 2, pp. 660–680, 2017. doi: 10.1007/s00453-016-0187-y.

[2] P. S. Oliveto and C. Zarges, "Analysis of diversity mechanisms for optimisation in dynamic environments with low frequencies of change," *Theoretical Computer Science*, vol. 561, part A, pp. 37–56, 2015. doi: 10.1016/j.tcs.2014.10.028.

[3] J.-B. Mouret and J. Clune, "Illuminating Search Spaces by Mapping Elites," arXiv:1504.04909, 2015.

[4] J. K. Pugh, L. B. Soros, and K. O. Stanley, "Quality Diversity: A New Frontier for Evolutionary Computation," *Frontiers in Robotics and AI*, vol. 3, art. 40, 2016. doi: 10.3389/frobt.2016.00040.

[5] J. Bossek and D. Sudholt, "Runtime Analysis of Quality Diversity Algorithms," *Algorithmica*, vol. 86, pp. 3252–3283, 2024. doi: 10.1007/s00453-024-01254-z.

[6] C. Qian, K. Xue, and R.-J. Wang, "Quality-Diversity Algorithms Can Provably Be Helpful for Optimization," arXiv:2401.10539, 2024.

[7] M. Zinkevich, "Online Convex Programming and Generalized Infinitesimal Gradient Ascent," in *Proceedings of the 20th International Conference on Machine Learning*, pp. 928–936, 2003.

[8] H. Lakkaraju, J. Kleinberg, J. Leskovec, J. Ludwig, and S. Mullainathan, "The Selective Labels Problem: Evaluating Algorithmic Predictions in the Presence of Unobservables," in *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2017. doi: 10.1145/3097983.3098066.

[9] C. F. Manski, "Partial identification with missing data: concepts and findings," *International Journal of Approximate Reasoning*, vol. 39, no. 2–3, pp. 151–165, 2005. doi: 10.1016/j.ijar.2004.10.006.

[10] A. Swaminathan and T. Joachims, "Counterfactual Risk Minimization: Learning from Logged Bandit Feedback," in *Proceedings of the 32nd International Conference on Machine Learning*, PMLR 37, pp. 814–823, 2015.

[11] Y. Jin, "Surrogate-assisted evolutionary computation: Recent advances and future challenges," *Swarm and Evolutionary Computation*, vol. 1, no. 2, pp. 61–70, 2011. doi: 10.1016/j.swevo.2011.05.001.